

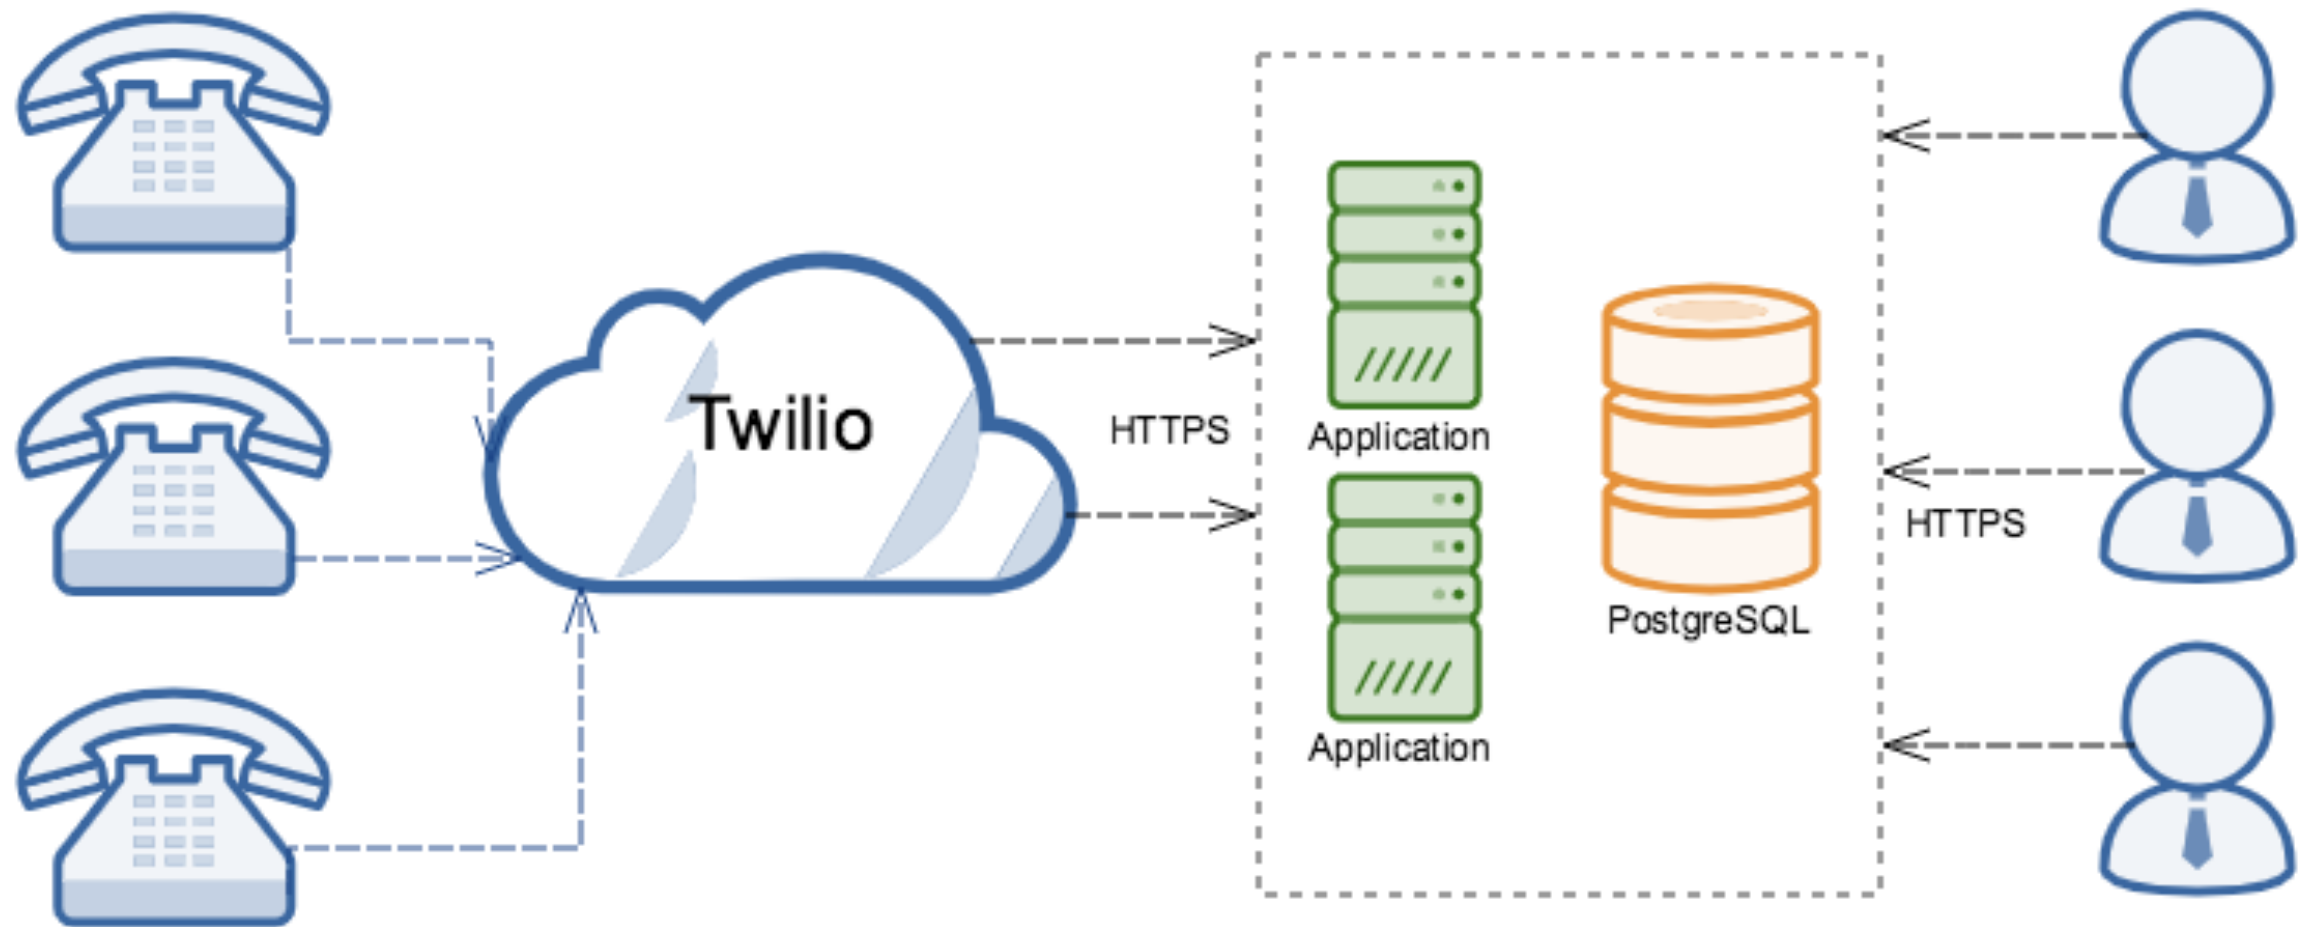
# Temporal data & time travel in PostgreSQL

Alexei Mikhailov, Glia

# Outline

- **How we usually investigate issues in production**
- **Temporal tables**
- **PostgreSQL extension `temporal_tables`**
- **Problems**
- **Alternative solutions**

# Call center



# Clients complaints

- Why was this call rejected?
- Why was this call picked up by operator X instead of Y?
- Why operator X had significantly less calls than others on day Y?

# Logs

# Logs

13:48:41.219	info	Operator updated	a925e5a5-52bd-43a4-8009-55fdf65b6a24	temporarily-unavailable
13:48:51.542	info	Operator removed	a925e5a5-52bd-43a4-8009-55fdf65b6a24	temporarily-unavailable
13:48:59.280	info	Operator added	b29beeaf-2863-4d27-984b-b6b8b3ac5465	available
13:48:59.298	info	Queue availability statuses are updated	-	-
13:49:12.172	info	Operator updated	b29beeaf-2863-4d27-984b-b6b8b3ac5465	temporarily-unavailable

# Logs

- **Incomplete**

> 2019-03-10 11:11 Operator is updated  
*What operator? What was updated?*

- **Personally identifiable information (PII)**

- **Reactive**

*Didn't have logs back then? GL HF*

- **Hard to use**

> 2018-03-10 23:00 Operator is ready, id=007  
> 2018-03-12 01:00 Operator is offline, id=007

*Search logs for date 2018-11-11: no results for agent 007*

- **Not durable \***

# Temporal tables

- **Debugging**
- **Audit and security**
- **Undo functionality**
- **BI**



# Temporal tables

- Big new feature of SQL 2011  
[https://cs.ulb.ac.be/public/\\_media/teaching/infoh415/tempfeaturessql2011.pdf](https://cs.ulb.ac.be/public/_media/teaching/infoh415/tempfeaturessql2011.pdf)
- Introduces new syntax  
... **FROM** table **FOR SYSTEM\_TIME AS OF** '<timestamp>'

\$\$\$ Oracle

Yes Since 10g+ (2005)

\$\$\$ IBM DB2

Yes Since version 10 (2010)

\$\$\$ MS SQL Server

Yes Since SQL Server 2016

PostgreSQL

No 3rd party extension temporal\_tables

MySQL

No

# Temporal tables

- Tables have additional time-period (interval) information
- Time period is closed-open (inclusive from the left, exclusive from the right)
  - PostgreSQL: `tstzrange( ' [2018-01-01, 2018-12-01) ' )`
- **System time-period:** what was the ACID state of my data on <time>?
  - Maintained by system
  - **UPDATE / DELETE / INSERT** automatically maintain this info
  - Only supports history, must never be future
  - Used by `...FROM table FOR SYSTEM_TIME AS OF`
  - **Usually separate table**

# Temporal tables

- **Business time-period**

- Maintained by application / user
- Represent business reality, e.g. "This contract will have price X for next 3 months"

contract_id	price	valid_period
1	50	[2018-01-01, 2018-02-01)
1	100	[2018-02-01, 2019-02-01)

- Can be in the future
- No special syntax for SELECT
- **Bi-temporal tables** = tables with system period and business period

# Historical table

```
CREATE TABLE agents (  
    id integer NOT NULL,  
    status character varying NOT NULL,  
    system_period tstzrange NOT NULL  
);
```

```
CREATE TABLE agents_history (LIKE agents);
```

# Historical table

```
CREATE FUNCTION save_previous_version() RETURNS TRIGGER AS $body$
BEGIN
    IF (TG_OP = 'INSERT') THEN
        NEW.system_period := tstzrange(current_timestamp, NULL, '[]');
        RETURN NEW;
    ELIF (TG_OP = 'UPDATE') THEN
        NEW.system_period := tstzrange(current_timestamp, NULL, '[]');
        INSERT INTO agents_history (id, status, system_period) VALUES
        (OLD.id, OLD.status, tstzrange(lower(OLD.system_period),
        current_timestamp, '[]'));
        RETURN NEW;
    ELSE
        INSERT INTO agents_history (id, status, system_period) VALUES
        (OLD.id, OLD.status, tstzrange(lower(OLD.system_period),
        current_timestamp, '[]'));
        RETURN OLD;
    END IF;
END;
$body$
LANGUAGE plpgsql;
```

```
CREATE TRIGGER history_trigger
BEFORE INSERT OR UPDATE OR DELETE ON agents
FOR EACH ROW EXECUTE PROCEDURE save_previous_version();
```

# Historical table

```
INSERT INTO agents (id, status) VALUES (1, 'ready');
```

## "agents"

id	status	system_period
1	ready	["18:59:17",)

(1 row)

## "agents\_history"

id	status	system_period
----	--------	---------------

(0 rows)

# History table

```
UPDATE agents SET status = 'away' WHERE id = 1;
```

## "agents"

id	status	system_period
1	away	["19:10:47",)

(1 row)

## "agents\_history"

id	status	system_period
1	ready	["18:59:17","19:10:47")

(1 row)

# History table

```
UPDATE agents SET status = 'offline' WHERE id = 1;
```

## "agents"

id	status	system_period
1	offline	["19:15:06",)

(1 row)

## "agents\_history"

id	status	system_period
1	ready	["18:59:17", "19:10:47")
1	away	["19:10:47", "19:15:06")

(2 rows)



# History table

```
DELETE FROM agents WHERE id = 1;
```

**"agents"**

id	status	priority	valid_from	valid_to
(0 rows)				

**"agents\_history"**

id	status	system_period
1	ready	["18:59:17", "19:10:47"]
1	away	["19:10:47", "19:15:06"]
1	offline	["19:15:06", "19:17:47"]
(3 rows)		

# Queries

```
CREATE VIEW agents_with_history AS  
SELECT * FROM agents  
UNION ALL  
SELECT * FROM agents_history
```

# Queries

**How did the data look like at point in time T?**

```
SELECT * FROM agents_with_history
WHERE id = 1
      AND '19:12'::timestamptz <@ system_period
```

id	status	system_period
1	away	["19:10:47","19:15:06")

(1 row)

```
SELECT * FROM agents_with_history
WHERE id = 1
      AND '19:09'::timestamptz <@ system_period
```

id	status	system_period
1	ready	["18:59:17","19:10:47")

(1 row)

# Queries

How did the data change in interval (T1, T2)

```
SELECT * FROM agents_with_history
WHERE id = 1
      AND tstzrange(' [18:55,19:12]') <@ system_period
```

id	status	system_period
1	ready	["18:59:17","19:10:47")
1	away	["19:10:47","19:15:06")

(2 rows)

```
SELECT * FROM agents_with_history
WHERE id = 1
      AND tstzrange(' [19:05,)') <@ system_period
```

id	status	system_period
1	away	["19:10:47","19:15:06")
1	offline	["19:15:06","19:15:47")

(2 rows)

# Extension

- [https://github.com/arkhipov/temporal\\_tables](https://github.com/arkhipov/temporal_tables)
- `pgxnclient install temporal_tables`
- `CREATE EXTENSION temporal_tables`  
*Requires superuser*
- `CREATE TRIGGER versioning_trigger  
BEFORE INSERT OR UPDATE OR DELETE ON agents  
FOR EACH ROW EXECUTE PROCEDURE  
versioning ('system_period', 'agents_history', true);`
- Doesn't require same set of columns

# Advantages

- **ACID**
- **Backups**
- **Easy to integrate with existing database**
- **Low performance overhead**

# Data duplication

- **Normalization**

*Move heavy columns out of the main table*

***Before:*** *users (id, name, picture)*

***After:*** *users(id, name, picture\_id), pictures(id, data)*

- **Remove heavy columns from historical table**

*Data will be lost*

- **Prune historical tables**

- **Separate tablespace on cheap HD**

# Caveats

- **Structure is not in sync with main table**
  - Can use inheritance, but with caution
  - Event triggers (i.e. *do something when schema changes*)
- **No primary key, no foreign keys**
- **You MUST use tstzrange for system\_period column**  
*If you're using temporal\_tables extension. Only "gist" index is supported.*
- **"Custom" extension**



# Alternative solutions

- **RDBMS with support for temporal tables**

*Oracle, IBM DB2, MS SQL*

- **Event sourcing**

- **Immutable data**

agents		statuses
+-----+		+-----+
id	<--	agent_id
name		status
		timestamp

- **Application-managed audit logs**

# Demo

Site ID(s)

cc0a1ccc-e00a-4b28-9ea5-277548995fec x

Queue ID(s)


Point in time

dd/mm/yyyy, --:--:--

(GMT-06:00) Central Time (US ↕)

Show

Queue Name	ID	Status	Enqueued visitors	Staffing operators
General	8788b207-edf8-4a66-a27f-ccd73b565e60	unstaffed	0	0



# Demo

Site ID(s)

cc0a1ccc-e00a-4b28-9ea5-277548995fec ✕

Queue ID(s)

Point in time

dd/mm/yyyy, --:--:--

(GMT-06:00) Central Time (US ↕

Show

Queue Name	ID	Status	Enqueued visitors	Staffing operators
General	8788b207-edf8-4a66-a27f-ccd73b565e60	unstaffed	0	0

# Recap

- Logs are not great
- Temporal tables are possible in any DB
- Use with caution
- Know alternatives

# Questions?

Slides:

<https://github.com/take-five/temporal-tables-presentation>

